

CS354 Midterm 2

Posix brk & unistd.h - contains collection of Posix API functions

Portable OS interface: standard for maintaining compatibility among Unix OS's

DIY Heap using Posix calls.

brk: points to end of program (shifts up if need more space). points at top of Heap.

① `int brk(void *addr)` : sets top of heap to specified addr (0 if successful, else -1 & sets errno).
- not common to use.

② `void *sbrk(intptr_t incr)` `intptr_t` is size of long for ptr addr.
- attempts to change program's top of heap by `incr` bytes.
- returns old brk if successful, else -1 & sets errno.

③ `errno` `#include <errno.h>`. `strerror(errno)` to convert to string, only 1 error at a time, overwritten if new error.

④ best to use `malloc/calloc/realloc/free`, since well-tested, cannot use combination \Rightarrow UB undefined program behaviour.

stdlib.h: 25 commonly used functions.

`int` 4 bytes long &.

- conversion: `atoi`, `stol` (string to long)
- execution flow: `skip to exit`, `exit(1)`
- math funcs: `absolute`, `floor`, `ceiling`
- 'search': `linear`, `binary`
- sorting: `qsort`
- random number: `rand`.

Heap Allocator Functions

`void* malloc(size_t size)` - allocates `size` (in bytes), returns pointer to block of heap memory, NULL if fail.

`void* calloc(size_t nItems, size_t size)` - Allocates, clears to 0, returns block of memory of `nItems * size` bytes, return NULL if allocation fails.

`void* realloc(void* ptr, size_t size)` - reallocates to `size` bytes a prev allocated block of heap mem pointed to by `ptr`, return NULL if fail.

• if `size` smaller than current \Rightarrow no problem, else make block bigger w/o interfering w other blocks (uses `malloc`)

`void free(void* ptr)` - free heap mem pointed to by `ptr`, does nothing if `ptr` is NULL.

Heap

- segment of Virtual Address Space used for dynamically allocated memory → requested while program is running.

- is a collection of various sized memory blocks.

↳ block: contains payload, overhead, allocator.

↳ payload: part requested & used by user

↳ overhead: part used by allocator to manage allocation

↳ allocator: code that allocates & frees.

- C uses explicit approach: tell how many bytes, explicitly call free to free up memory
- vs Java "new" calculates bytes needed, "garbage collector", free mem when needed.

- Allocator Design Goals.

1. Maximize Throughput - request bytes in certain order, just going to allocate the space } Trade off: improve one, worsens other.
↳ # of malloc & frees per time interval, higher better.
2. Maximize Memory Utilization - % of memory used, that is used for payload.
↳ memory requested (bytes) / heap that has been allocated (bytes).

- Double Word Alignment

① Block sizes multiple of 8 ② Payload addr at multiple of 8.

- Fragmentation: ① External - enough free heap mem divided into blocks that are each too small.
- happens due to sequence of calls.


② Internal - allocated memory used for overhead instead of payload.
- block padding, block headers causes this.

③ Fake - when combined memory in adjacent free blocks sufficient to satisfy request but individually not.

- Explicit Free List (EFL): allocator maintains data structure with free blocks.

-ve: may need more memory for data structs.

fre: time to search only free block for size.



- Implicit Free List: loop all blocks instead of just list of free blocks, allocator uses heap blocks as data structure.
- code must track size & status of each block.

-ve: time, more time to search both alloc'd & free blocks

fre: may be less memory.

- Placement Policies: Assuming heap pre-divided into various size free blocks smallest to largest

① First Fit: start from beginning, stop at first free block that is big enough, fail if reach END MARK.

mem util: low, likely choose desired size throughput: -ve, must step through many allocated blocks.

② Next Fit: start from block most recently alloc'd, stop at first free block that is big enough, fail if reach block started with (wrap around)
mem util: not as good, may choose too big a block throughput: faster, skip many prev alloc'd blocks.

③ Best Fit: start from beginning, stop END MARK & choose best fit or early if EXACT FIT, fail if no block big enough.
mem util: best, closest to size of request throughput: slowest in general.

- Free Block

- If bigger than request

① Use entire block

↳ memutil: -ve, ↑ internal fragmentation — waste space.

↳ thput: +ve, faster code.

② Split into used + free blocks

↳ memutil: +ve, less wasted space.

↳ thput: -ve, more operations, slower to search, more blocks to search.

- If request too large for any free block.

① coalesce adjacent free blocks. (cannot move already allocated blocks to create larger free areas).

② ask kernel for more.

③ return NULL, alloc fail.

- Coalescing Free Blocks. (constant time).

- immediate: coalesce after freeing block. — check prev & next, coalesce if possible.

- delayed: only if needed to satisfy request — in alloc func

- Free Block Footers

• last word of free block (4 bytes) contain free block size, allocated blocks no need since cannot be freed.

• from payload to prev block HDR: $(ptr-4) \Rightarrow$ check p-bit, if 0, $(ptr-4) - prev_block_size$

* take note scaling (void*) or (char*) for 1.

- Explicit Free List Layout



* Footer still useful for faster coalescing.

* order of free blocks in free list don't have to be in same order as found in addr space.

Improvements: — Free List Ordering

① address order: maintain list in order from low to higher addr.

malloc with FF: easier to find free block. +ve

free: slower -ve

② last-in-order: append most recently freed block at end of list.

malloc with FF: slower, must go through most recent

free: faster, not trying maintain order.

— Free List Segregation

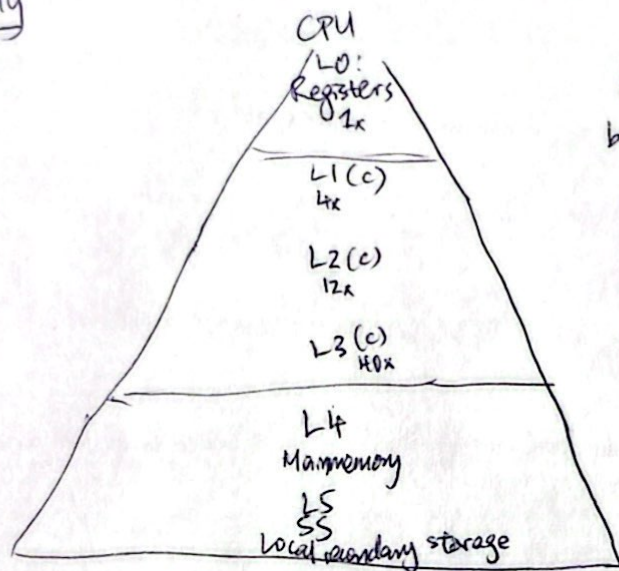
① simple segregation: one EFL for each block size

② fitted segregation: one EFL for each size range — small, medium, large.

+ve memutil — as good as best fit

+ve thput — search only part of the heap.

Memory Hierarchy



word: size btwn L1 & CPU

block: transfer btwn C levels & MM

page: sized used by MM

trf between MM & SS.

• cpu cycles: used to measure time

• latency: memory access time (delay)

Temporal Locality: recently used location repeatedly used in future.

Spatial Locality: recently used location is followed by access of nearby location. (on the same pallet).

Stride: step size in words between sequential access in the block. Smaller better.

Cache Block: unit of transfer between MM & C levels, typically 32 bytes.

Good or Bad Locality?

Instruction Flow:

- sequencing: good spatial

- selection: if...else... bad spatial, bad temporal.

- repetition: good temporal, good spatial if stride is small.

Searching Algorithms:

- Linear search: good spatial if linear array, bad array if linked list/tree } some temporal.

- Binary search: bad spatial

Cache Misses: ^{block not in cache} ① Cold miss: ~~cache~~ location empty ② Capacity Miss: cache too small for working set ③ Conflict miss: 2 or more blocks map to same location.

Cache Hit: Block in cache.

Placement Policies: ① unrestricted - can go anywhere (L1) ② Restricted - mapped to set of locations (L2). usually block num % 16.

Replacement Policies: ① Choose any block location (L1) ② Restricted - block # % 16 (L2).

Victim Block: Block to be replaced

Working set: set of blocks accessed given some interval of time.

Bytes in Addr Space

Let $M = \# \text{bytes}$. (IA-32 4GB).

so $M = 2^m \Rightarrow m = \log_2 M = \# \text{bits in addr (32 in IA-32)}$.

How Big is a block?

- Cache blocks should be big enough to capture spatial locality, small enough to minimize latency.

Let $B = \# \text{bytes per block}$. (IA-32, 32 bytes per block).

$B = 2^b = 32 \text{ bytes} \Rightarrow b = \log_2 B = 5$ (5 b bits).

b-bits: #bits of address to determine which byte of block.

word offset: identify which word of block.

byte offset: identify which byte of word (last 2 bits). since 4 byte/word.

b-bits on most right. If in most left side \Rightarrow LOSE SPATIAL LOCALITY.

• all bytes near each other won't be in same block.

How many 32-byte blocks in 32-bit addr space?

$$MAS/B = 2^{32}/2^5 = 2^{27} = 2^7 2^{20} = 128 \text{ Mb.}$$

For 32 byte blocks, if last 5 bits are $\begin{array}{c} 01101 \\ \hline \text{3rd} \quad 1 \end{array}$ use byte 1 of word 3.

Sets & Tags

Set: where a block of memory is uniquely mapped, maps to specific set in cache.

tag: uniquely identifies particular block in set.

$S = \# \text{sets in cache} = 2^s \Rightarrow s = \# \text{bits for sets}$.

s-bits \Rightarrow identifies which set block maps to.


why s bits not most left \Rightarrow don't want consecutive blocks in same set, might displace each other will lose spatial locality.

How many blocks map to each set for 32-bit AS and a cache with 1024 sets?

$$(MAS/S)/B = (2^{32}/2^{10})/2^5 = 2^7 = 2^7 2^{10} = 128 \text{ Kbytes}$$

t-bits: bits to identify which block in set. * when block is copied into cache, t bits stored as well

v-bit: indicate if cache line is in use or not



Request for word

1. Set Selection: using s bits (gets set)

2. Line Matching: using t-bits match tag, v-bit must be valid too. (gets line).

For L1 cache only (highest level), must extract word from Block as well

Direct Mapped Cache: cache with s sets, 1 line per set.
mem blocks map to exactly 1 set.

cache operations: • no search, set selection $O(1)$, line matching $O(1)$
• simple circuitry.

when 2 mem blocks map to same set \Rightarrow CONFLICT MISS, same s bits, diff t bits so diff block num
but since only 1 line/set, only 1 block can be stored.

* can cause thrashing: continuously exchange blocks.

Appropriate for larger caches.

Fully Associative Cache: cache with 1 set, E lines in that set.

mem blocks can be at any line in that set.

cache operations: • set selection $O(1)$, line matching $O(E)$ where E is # of lines.
• complex circuitry since every block matched to every line, many things to match.

when 2 diff blocks map to same set \Rightarrow all map to same set \Rightarrow choose free line.

How many lines? $C = S \times E \times B$, $S=1$, so $E = \frac{C}{B}$.

Appropriate for smaller caches (L1).

Set Associative Cache: S sets with E lines per set.

mem blocks map to a single set with E lines (can be any line in that set).

cache operations: set selection $O(1)$, line matching $O(E)$ E smaller than Fully Assoc.
compute from s bits

+ve: \downarrow conflict miss -ve: more circuitry than direct mapped
less circuitry than fully assoc.

$E=4$
4-way associative

$C = (S, E, B, m)$ ^{# bits in addr}

$C = S \times E \times B$ For $(1024, 4, 32, 32)$

$C = 2^{10} \times 2^2 \times 2^5 = 2^{17}$ with 15 + bits.

Replacement Policies

- ① Random Replacement: last block must be in cache, usually no duplicates & empty lines.
- ② Least Recently Used: use queue keep track.
- ③ Least Frequently Used: track how often each line is used, each line has a count.
 - set to 0 if line gets new block.
 - increment when line is accessed
 - if tie, evict randomly chosen of ties.

Writing to a cache

- Reading data copies a block of memory into cache levels.
- Writing data requires these copies are consistent.

Write Hits: writing to a block that is in the cache (K)

1. Write Through: immediate update lower level (K+1)
 - ve: need wait for lower level to write, ↑ time
 - ve: more bus traffic, passing all writes through
 - +ve: or add bypass cache with write buffer (smaller size cache to let it wait for K+1 to do write).
2. Write Back: write to next level (K+1) when block evicted from K (line changes)
 - +ve: faster, no wait
 - ve: less bus traffic, only update when done w block.
 - ve: must track if line has been changed add "DIRTY" bit

Write Misses: when writing to a block that is not in cache (K).

1. No Write Allocate: write directly to next lower level (K+1) bypass cache K.
 - ve: must wait for next lower level to write (K+1)
 - +ve: less bus traffic since no wait to get block.
2. Write Allocate: read into curr level K first, then write
 - ve: must wait for block to be read in
 - ve: more bus traffic, must pass block to current cache level.

Typical Designs: ① Write Through with No Write Allocate "get to K+1"

② Write Back with Write Allocate "keep at K"

③ exploits locality test, keeps in highest level before having to go lower level.

Cache Performance

Metrics: ① Hit rate: # hits / # mem access, higher better.

② Hit time: time to determine cache hit, set selection + line matching, lower better.

③ Miss Penalty: additional time to process a miss, lower better.

Larger Blocks (S & E unchanged) $C = S \times E \times B$

- hit rate: better, more spatial locality per block, eg more values from array in same block.
- hit time: same, since set selection & line matching same
- miss penalty: worse, ↑ bus traffic when transferring block.
- Thus block sizes small, usually 32/64 bytes.

More Sets (E & B unchanged)

- hit rate: better, ↑ temporal locality, more sets so diff blocks can be in cache at the same time.
- hit time: worse, more sets, more complexity (circuitry) to find correct matching set (still $O(1)$)
- miss penalty: same, since block size same
- Thus, faster caches smaller (L1), slower caches larger.

More Lines E per set (S & B unchanged)

- hit rate: better, more temporal locality, more lines, more ^{diff} blocks in cache at same time, fewer conflict misses.
- hit time: worse, slower line matching $O(E)$
- miss penalty: ~~same~~ worse, more lines, longer to detect match. Therefore, faster caches fewer lines (L1)