

CS 354 - Machine Organization & Programming

Tuesday Feb 13th and Thursday Feb 15th, 2024

Midterm Exam - Thursday, February 22nd, 7:30 - 9:30 pm

- ◆ Room: Students will be assigned a room and sent email with that room
- ◆ UW ID required
- ◆ #2 pencils required
- ◆ closed book, no notes, no electronic devices (e.g., calculators, phones, watches)
- ◆ see "Midterm Exam 1" on course site Assignments for topics

PM BYOL: Start p2A and p2B if you have not yet started either

Activity A04: due on or before this week Saturday

Homework hw1: Due on or before this week Monday (solution available Wed morning)

Homework hw2: Due on or before next week Monday

Project p2A: Due on or before this week Friday, Feb 16

Project p2B: Due on or before next week Friday, Feb 23

Week 4 Learning Objectives (at a minimum be able to)

- ◆ use `<stdio.h>` functions: `printf`, `scanf`, `fopen`, `fclose`, `fgets`, `fputs`
- ◆ use predefined file pointers: `stdin` and `stdout`
- ◆ use format specifiers: `%c` `%f` `%i` `%d` `%s` `%p` `%x`
- ◆ use Linux I/O redirection at the command line: `< input_file > output_file >> append_file`
- ◆ describe C's abstract memory model: **Process View = Virtual Memory**
- ◆ diagram C's abstract memory model: **CODE, DATA, HEAP, STACK**
- ◆ meet IA-32 memory hierarchy: **Hardware View = Physical Memory**
- ◆ understand difference and use of **global** vs **static local** variables

This Week

Pointers to Structures (from last week) Standard & String I/O in <code>stdio.h</code> File I/O in <code>stdio.h</code> Copying Text Files Three Faces of Memory Virtual Address Space	C's Abstract Memory Model Meet Globals and Static Locals Where Do I Live? Linux: Processes and Address Spaces Exam Sample Cover Page
<p>Next Week: The Heap & Dynamic Memory Allocators (p3)</p> <p>Read: B&O 9.1, 9.2, 9.9.1-9.9.6</p> <p>9.1 Physical and Virtual Addressing 9.2 Address Spaces 9.9 Dynamic Memory Allocation 9.9.1-9.9.6</p>	

Standard and String I/O in stdio.h

Standard I/O

string: sequence of char w null terminating char

Standard Input

getchar //reads 1 char gets char from std input by default, reads 1 char and leaves rest on stream
gets //reads 1 string ending with a newline char, BUFFER MIGHT OVERFLOW
string: sequence of char w null terminating char

int scanf(const char *format_string, &v1, &v2, ...)

reads formatted input from the console keyboard
returns number of inputs stored, or EOF if error/end-of-file occurs before any inputs
End of File

format string contains format specifiers and characters to skip

format specifiers %d, %f float, %p, %s string, %i int 03 octal

must match corresponding dest variable 0x1b hex
33 dec

%d: formats the output as a signed decimal integer.

whitespace input separator (space, tab, newline)

Standard Output

putchar //writes 1 char
puts //writes 1 string

int printf(const char *format_string, v1, v2, ...)

writes formatted output to the console terminal window
returns number of characters written, or a negative if error

format string format specifiers and chars to display

use \n to flush buffer

Standard Error alternate output location, terminal window by default

void perror(const char *str)

writes formatted error output to the console terminal window

can read from and write to string instead of keyboard terminal

String I/O

scan from string, into the variables
int sscanf(const char *str, const char *format_string, &v1, &v2, ...)
reads formatted input from the specified str
returns number of characters read, or a negative if error

int sprintf(char *str, const char *format_string, v1, v2, ...)
writes formatted output to the specified str
returns number of characters written, or a negative if error

File I/O in stdio.h

Standard I/O Redirection in Linux

a.out (executable, presumably requires input)	a.out (executable, presumably has output)
a.out < input_file	a.out > output_file overwrites to output_file
	a.out >> output_file appends to output_file

File I/O

File Input

```
fgetc/getc, ungetc //reads 1 char at a time  
fgets                //reads 1 string terminate with a newline char or EOF  
                      expects data file to be opened  
int fscanf(FILE *stream, const char *format_string, &v1, &v2, ...)  
                      reads formatted input from the specified stream  
                      returns number of inputs stored, or EOF if error/end-of-file occurs before any inputs
```

File Output

```
same as about but to file instead of terminal keyboard now  
fputc/putc          //writes 1 char at a time  
  
fputs                //writes 1 string  
int fprintf(FILE *stream, const char *format_string, v1, v2, ...)  
                      writes formatted output to the specified stream  
                      returns number of characters written, or a negative if error
```

Predefined File Pointers

```
stdin    is console keyboard  
stdout    is console terminal window  
stderr    is console terminal window, second stream for errors
```

“r” read or “w” write

Opening and Closing

```
FILE *fopen(const char *filename, const char *mode)  
                      opens the specified filename in the specified mode  
                      returns file pointer to the opened file's descriptor, or NULL if there's an access problem  
                      check for errors  
  
int fclose(FILE *stream)  
                      flushes the output buffer and then closes the specified stream  
                      returns 0, or EOF if error  
  
                      error if no permission to write/read etc
```

Copying Text Files

name of file is copy

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {

    if (argc != 3) {
        fprintf(stderr, "Usage: copy inputfile outputfile\n");
        exit(1);
    }

    FILE *ifp = fopen(argv[1], "r");
    if (ifp == NULL) {
        fprintf(stderr, "Can't open input file %s!\n", argv[1]);
        exit(1);
    }

    FILE *ofp = fopen(argv[2], "w");
    if (ofp == NULL) {
        fprintf(stderr, "Can't open output file %s!\n", argv[2]);
        fclose(ifp);
        exit(1);
    }

    const int bufsize = 257; //WARNING: assumes lines <= 256 chars
    char buffer[bufsize];
    while ( fgets(buffer, bufsize, ifp) != NULL ) {
        fputs(buffer, ofp);
    }

    if (fclose(ifp) == NULL) {
        print("error ifp close\n");
        exit(1);
    }

    if (fclose(ofp) == NULL)
        print("error ofp close\n");
        exit(1);
}

}
```

copy src_filename dst_filename

Three Faces of Memory

* **Abstraction:** manage complexity by focusing on relevant details

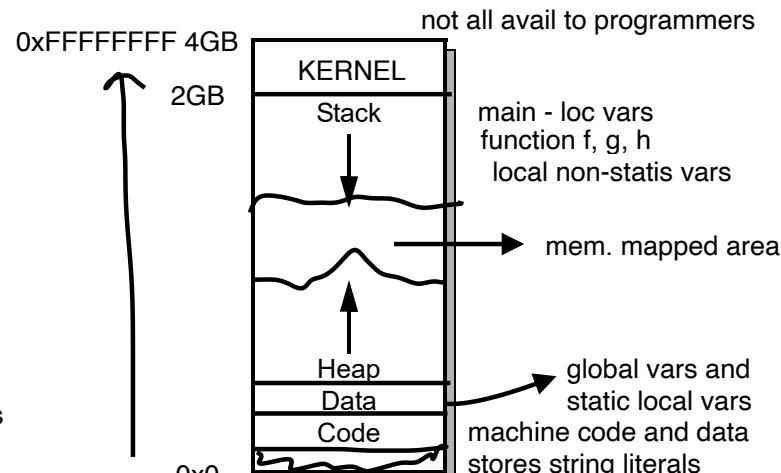
Process View = Virtual Memory

Goal: provide simple view

virtual address space (VAS):

illusion by OS that each process has its own contiguous memory block

virtual address: simulated addr a process generates



System View = Illusionist (CS 537)

Goal: make memory sharable and secure

pages: 4KB units of storage

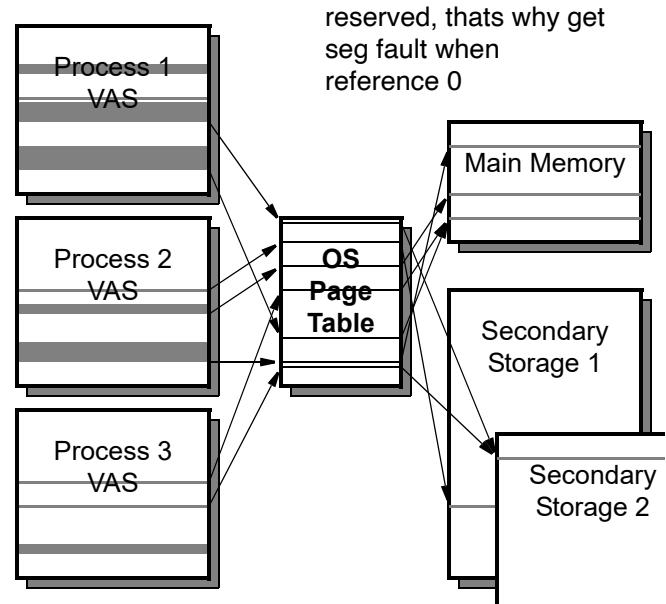
fixed-size determined by OS, contiguous

page table:

OS data structure

maps Virtual pages to Physical pages

keeps process from interfering with other processes



Hardware View = Physical Memory

Goal: keep CPU busy

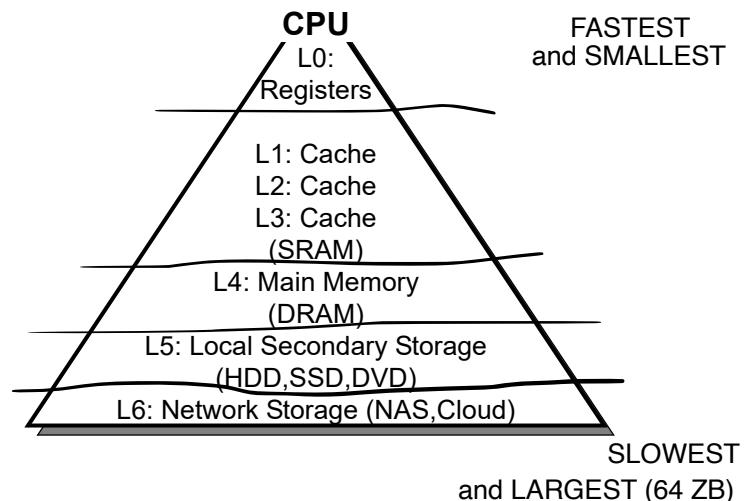
physical address space (PAS):

multi-level hierarchy

frequently accessed data in memory closest to CPU

physical address:

actual address used to access memory



PROCESS VIEW
Virtual Address Space (IA-32/Linux)

32-bit Processor = 32-bit Addresses => $2^{32} = 4,294,967,296 = 4\text{GB}$ Address Space

4GB = 11111111111111111111111111111111 = 0xFFFFFFFF

address space:

range of valid memory addresses for processes

process: a running program

Windows IA-32 has 2GB of OS, user gets 2GB

11000000000000000000000000000000 = 0xC0000000

r/w

kernel: memory resident portion of the OS

user process: process that is not kernel

* *Every user process* has this simple view of memory
makes coding easier

Dynamic Linked Libraries (DLL), File I/O, large memory requests ← r/w

magic number - where code segment starts



00001000000010010000000000000000 = 0x08048000

00000000000000000000000000000000 = 0x00000000

binary addresses

hex addresses

brk
break pointer

r/w

r/w

r/o

section

globals
static local

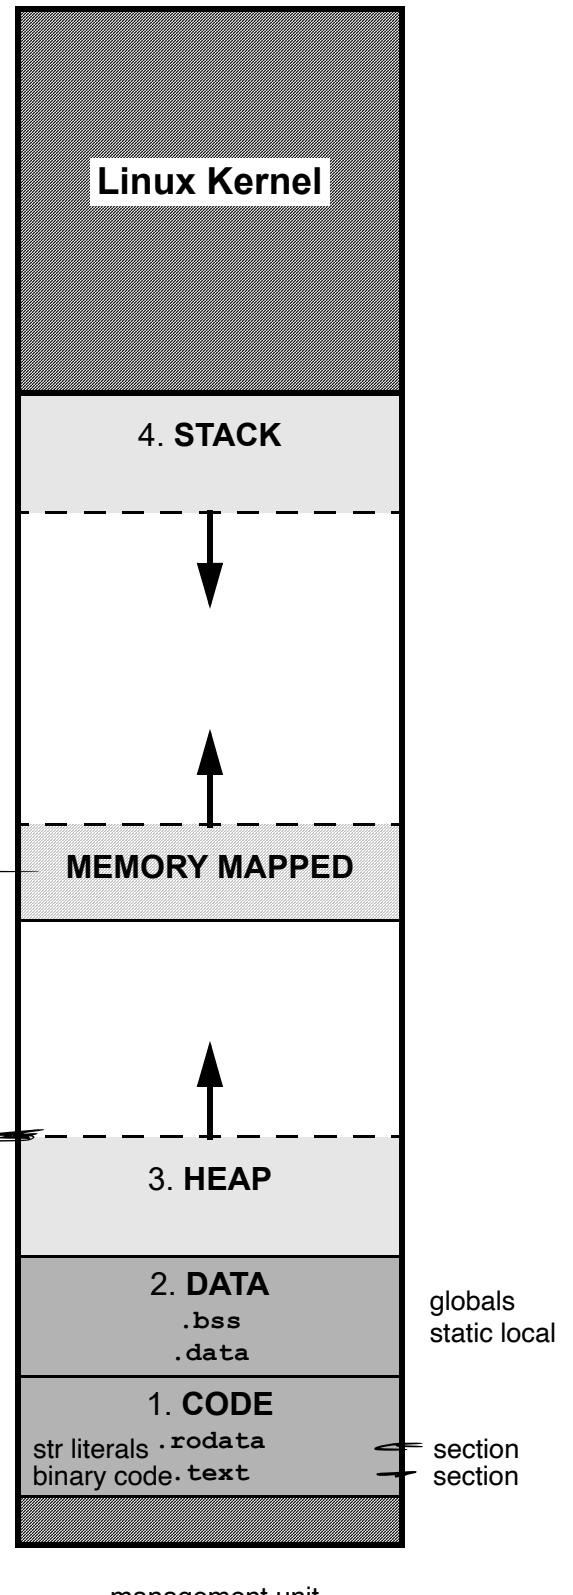
3. HEAP

2. DATA
.bss
.data

1. CODE

str literals
.rodata
binary code
.text

management unit



C's Abstract Memory Model

1. CODE Segment

Contains: program code — binary machine code

.text section machine code

.rodata section str literals

Lifetime: entire program's execution

LOADER put code in memory before start executing

Initialization: by LOADER from executable object file (before execution begins)

Access: READ ONLY access

2. DATA Segment

Contains: global variables and static locals

Lifetime: entire program's execution

Initialization: by LOADER from executable object file before execution

.data section variables that are initialized to non-zero values

.bss section variables that are not initialized or initialized to zero

Block started by symbol

Access: read/write

3. HEAP (AKA Free Store)

Contains: memory allocated and freed by the programmer during execution

Lifetime: managed by the programmer — malloc/calloc/realloc — free

can malloc in one func, can free in another func

Initialization: none by default

Access: read/write

4. STACK (AKA Auto Store)

Contains: memory in stack frames — automatically allocated and freed by code generated by compiler

stack frame (AKA activation record) non-static local , parameters, temp local vars (i,j,k)

int i = 11;
non-static

Lifetime: from when its declared till end of scope

Initialization: none by default

Access: read/write

Meet Globals and Static Locals

What?

A global variable is

- ◆ declared outside of any function
- ◆ accessible to all functions in the source file
- ◆ allocated in the DATA segment

A static local variable is

- ◆ declared inside a function with a static modifier
- ◆ accessible only within the function
- ◆ allocated in the DATA segment

Why?

for storage that exists during entire runtime

* *In general, global variables* should not be used!

Instead use local variables that are passed to callee functions

How?

```
#include <stdio.h>
int g = 11;    global variable
               non-static local
void f1(int p) {
    static int x = 22;  static local variable
    x = x + p * g;
    printf("%d\n", x);
}

int main(void) {
    f1(g);
    g = 2;
    int g = 1;    non-static local, shadowed global variable g
    f1(g);
    return 0;
}
```

Output
143
145

shadowing: when local variable blocks access to global variable

* *Avoid shadowing; don't use the same identifier* for local and global variables

Where do I live?

→ Identify the segment (and section) for each memory allocation in the code below.

```
#include <stdio.h>                                         highlighted is memory allocations
#include <stdlib.h>

int gus = 14;      // global variable, in DATA segment, .data file
int guy;          // global variable, in DATA segment, .bss file

int madison(int pam) {           // non-static local variable, in STACK segment
    static int max = 0;          // static local variable, in DATA segment, .bss file
    int meg[] = {22, 44, 88};    // non-static local variable, in STACK segment      SAA
    int *mel = &pam;            // non-static local variable, in STACK segment
    max = gus--;
    return max + meg[1] + *mel;
}

int *austin(int *pat) {          // non-static local variable, in STACK segment
    // (parameters treated like local variables non-static)
    static int amy = 33;         // static local variable, in DATA segment, .data file
    int *ari = malloc(sizeof(int)*44); // ari: // non-static local variable, in STACK segment
    gus--;                      memory ari points to is on heap
    *ari = *pat;
    return ari;
}

int main(int argc, char *argv[]) { // argc: // non-static local variable, in STACK segment
    // argv: // non-static local variable, in STACK segment
    int vic[] = {33, 66, 99};    // non-static local variable, in STACK segment
    int *wes = malloc(sizeof(int)); // wes: // non-static local variable, in STACK segment
    *wes = 55;                  memory wes points to is on heap
    guy = 66;
    free(wes);
    wes = vic;
    wes[1] = madison(guy);
    wes = austin(&gus);
    free(wes);                 string literals in CODE segment
    printf("Where do I live?\n");
    return 0;
}
```

* *Arrays, structs, and variables* can live in DATA, HEAP, or STACK segments

if array of strings, array still on stack but string literals on the code segment

pointers can store any address but can get SEG FAULT if try to dereference to address without permission

Linux: Processes and Address Spaces

program when running what can i do with program while running

Process and Job Control in reality not much

- ◆ Linux is multi-tasking OS, where you can run multiple processes concurrently

ps list snapshot of user processes, for everyone “ps -e”

jobs list only user processes, from command line

& put process in background

ctrl+z suspend running process

bg put suspended process in background

fg bring a process to foreground

ctrl+c stop running foreground process

top display table of resource usage processes

Program Size

size <executable or object_file> display size of program's memory segments
.rodata .text .data .bss

```
$gcc -m32 myProg.c
$size a.out
    text      data      bss      dec      hex filename
  1029        276         4     1309      51d a.out
  CODE      DATA
size in decimal and hex
```

Virtual Address Space Maps

- ◆ Linux enables you to see Virtual Address Space (mem map) of each process

\$pmap <pid_of_process>

\$cat /proc/<pid_of_process>/maps magic number, stack, libraries, Virtual Dynamic

\$cat /proc/self/maps notice heap

/proc: virtual file system that reveals KERNEL data in ASCII text from that can be read by progs

cat /proc/loadavg

SPEC CODES EF 10	UW LOGIN NAME	Last, First Name (as in email and on scantron)
------------------------	---------------	--

Computer Sciences 354
Midterm Exam 1 Primary

Thursday, October 5th, 2023

60 points (15% of final grade)

Instructor: Debra Deppele

1. RECORD Special Codes for EF on scantron. Ask Proctor if there are not 2 digits..
2. PRINT your UWNET ID (login name not photo id number) in box above.
3. PRINT Last, First Name in box above.
4. SCANTRON Fill in all fields and their bubbles on the scantron form
(must use #2 pencil on scantron form).
 - (a) LAST NAME field - left align last name as given in room email
 - (b) FIRST NAME field - left align first five letters of your first name as given in email
 - (c) IDENTIFICATION NUMBER your UW Student WiscCard ID number
 - (d) SPECIAL CODES E - write and fill-in bubble for your exam version number 1
 - (e) SPECIAL CODES F - write and fill-in bubble for your room number 0
5. FILL IN BUBBLES FOR ALL IDENTIFICATION FIELDS
and for SPECIAL CODES COLUMNS E and F.
6. Taking this exam indicates that you agree: to not write answers in large letters and to keep your answers covered; to not view or use another's work or any unauthorized devices in any way; to not make any type of copy of any portion of this exam; and that you understand that being caught doing any of these actions, or other actions that may permit any student to submit work that is not wholly their own will result in automatic failure of the exam and possible failure of the course. **Penalties are reported to the Deans Office for all involved.**

Parts	Number of Questions	Question Format	Possible Points
I	10	2 pt Simple Choice	20
II	12	3 pt Multiple Choice (+bonus)	36
III	4	Survey	4
	28	Total	60

Assumptions unless instructions explicitly state otherwise:

- + addresses and integers are 4 bytes unless explicitly stated otherwise.
- + code questions are about C std=gnu99 and IA-32 on our Linux platform

Reference: Powers of 2

$$2^5 = 32, 2^6 = 64, 2^7 = 128, 2^8 = 256, 2^9 = 512, 2^{10} = 1024 \\ 2^{10} = \text{K}, 2^{20} = \text{M}, 2^{30} = \text{G} \\ 2^A * 2^B = 2^{A+B}, 2^A / 2^B = 2^{A-B}$$

**Turn off and put away all notes and electronic devices and
wait for the proctor to signal the start of the exam.**